

Name: Saurabh Kumar Jha

Sec: B, Roll No. 53

LABSHEET 03

```
import pandas as pd
```

```
df = pd.read_csv("/workspaces/Mastery_in_DataAnalytics_and_Visualizations_and_Career_Advancemen/WA_Fn-UseC_-HR-Employee-Attrition-selected-columns.csv")
```

```
print(df.head())
```

```
print(df.shape)
```

```
print(df.columns.tolist())
```

	Age	Attrition	BusinessTravel	DailyRate	Department
0	41	Yes	Travel_Rarely	1102	Sales
1	49	No	Travel_Frequently	279	Research & Development
2	37	Yes	Travel_Rarely	1373	Research & Development
3	33	No	Travel_Frequently	1392	Research & Development
4	27	No	Travel_Rarely	591	Research & Development

	DistanceFromHome	Education	EducationField	EmployeeCount
0	1	2	Life Sciences	1
1	8	1	Life Sciences	1
2	2	2	Other	1
3	3	4	Life Sciences	1
4	2	1	Medical	1

```
(1470, 10)
```

```
['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',  
'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',  
'EmployeeNumber']
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from scipy import stats
```

```
import statsmodels.api as sm
```

```

print("Mean:")
print(df[["Age", "DailyRate", "DistanceFromHome"]].mean())

print("\nMedian:")
print(df[["Age", "DailyRate", "DistanceFromHome"]].median())

print("\nMode:")
print(df[["Age", "DailyRate", "DistanceFromHome"]].mode().iloc[0])

print("\nVariance:")
print(df[["Age", "DailyRate", "DistanceFromHome"]].var())

print("\nStandard Deviation:")
print(df[["Age", "DailyRate", "DistanceFromHome"]].std())

```

```

Mean:
Age                36.923810
DailyRate          802.485714
DistanceFromHome   9.192517
dtype: float64

```

```

Median:
Age                36.0
DailyRate          802.0
DistanceFromHome   7.0
dtype: float64

```

```

Mode:
Age                35
DailyRate          691
DistanceFromHome   2
Name: 0, dtype: int64

```

```

Variance:
Age                83.455049
DailyRate          162819.593737
DistanceFromHome   65.721251
dtype: float64

```

```

Standard Deviation:
Age                9.135373
DailyRate          403.509100
DistanceFromHome   8.106864
dtype: float64

```

```

print(df[["Age", "DailyRate", "DistanceFromHome"]].describe())

```

count	Age	DailyRate	DistanceFromHome
mean	1470.000000	1470.000000	1470.000000
std	36.923810	802.485714	9.192517
	9.135373	403.509100	8.106864

min	18.000000	102.000000	1.000000
25%	30.000000	465.000000	2.000000
50%	36.000000	802.000000	7.000000
75%	43.000000	1157.000000	14.000000
max	60.000000	1499.000000	29.000000

```
numeric_df = df[["Age", "DailyRate", "DistanceFromHome"]]
correlation = numeric_df.corr()
```

```
print(correlation)
```

	Age	DailyRate	DistanceFromHome
Age	1.000000	0.010661	-0.001686
DailyRate	0.010661	1.000000	-0.004985
DistanceFromHome	-0.001686	-0.004985	1.000000


```
plt.figure(figsize=(8, 6))
```

```
plt.imshow(correlation, cmap="coolwarm")
plt.colorbar()
```

```
plt.xticks(
    range(len(correlation.columns)),
    correlation.columns,
    rotation=45
)
```

```
plt.yticks(
    range(len(correlation.columns)),
    correlation.columns
)
```

```
plt.title("Correlation Matrix")
plt.tight_layout()
plt.show()
```



```
r, p = stats.pearsonr(df["Age"], df["DailyRate"])
```

```
print("Pearson Correlation:", r)
```

```
print("P-value:", p)
```

```
Pearson Correlation: 0.01066094264553816
```

```
P-value: 0.6829722453880914
```

```
left = df[df["Attrition"] == "Yes"]["Age"]
```

```
stayed = df[df["Attrition"] == "No"]["Age"]
```

```
print("Employees who left:", len(left))
```

```
print("Employees who stayed:", len(stayed))
```

```
Employees who left: 237
```

```
Employees who stayed: 1233
```

```
t_stat, p_value = stats.ttest_ind(
    left,
    stayed,
    equal_var=False
)
```

```
print("T-statistic:", t_stat) print("P-
value:", p_value)
```

```
T-statistic: -5.82801185398895 P-
value: 1.3797600649439802e-08
```

```
if p_value < 0.05:
    print("Reject H0")
    print("There is a significant difference in Age between employees
who left and stayed.")
else:
    print("Fail to reject H0")
    print("There is no significant difference in Age between employees
who left and stayed.")
```

```
Reject H0
```

```
There is a significant difference in Age between employees who left
and stayed.
```

```
print(df["Department"].unique())
```

```
<StringArray>
```

```
['Sales', 'Research & Development', 'Human Resources']
```

```
Length: 3, dtype: str
```

```
groups = []
```

```
for department in df["Department"].unique():
    groups.append(
        df[df["Department"] == department]["Age"]
    )
```

```
f_stat, p_value_anova = stats.f_oneway(*groups)
```

```
print("F-statistic:", f_stat)
```

```
print("P-value:", p_value_anova)
```

```
F-statistic: 0.7655031924976907
```

```
P-value: 0.4652855299935297
```

```
if p_value_anova < 0.05:
    print("Reject H0")
    print("There is a significant difference in Age among
Departments.")
else:
    print("Fail to reject H0")
```

```
print("There is no significant difference in Age among  
Departments.")
```

Fail to reject H0

There is no significant difference in Age among Departments.

```
X = df["Age"]  
y = df["DailyRate"]
```

```
X = sm.add_constant(X)
```

```
model = sm.OLS(y, X).fit()
```

```
print(model.summary())
```

OLS Regression Results

```
=====
```

Dep. Variable:	DailyRate	R-squared:	
0.000			
Model:	OLS	Adj. R-squared:	F-
-0.001			
Method:	Least Squares	statistic:	
0.1669			
Date:	Mon,31 Aug 2026	Prob (F-statistic):	Log-
0.683			
Time:	08:31:36	Likelihood: AIC:	
-10906.			
No. Observations:	1470	BIC:	
2.182e+04			
Df Residuals:	1468		
2.183e+04			
Df Model:	1		
CovarianceType:	nonrobust		

```
=====
```

```
=====
```

	coef	std err	t	P> t	[0.025
0.975]					

const	785.0985	43.847	17.905	0.000	699.089
871.108					
Age	0.4709	1.153	0.408	0.683	-1.790
2.732					

```
=====
```

```
=====
```

Omnibus:	1174.453	Durbin-Watson:	
2.082			

Prob(Omnibus):	0.000	Jarque-Bera (JB):
88.696		
Skew:	-0.004	Prob(JB):
5.49e-20		
Kurtosis:	1.797	Cond. No.
159.		

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
plt.scatter(df["Age"], df["DailyRate"])
```

```
plt.xlabel("Age")
```

```
plt.ylabel("Daily Rate")
```

```
plt.title("Age vs Daily Rate")
```

```
plt.show()
```



```
print("R-squared:", model.rsquared)
```

R-squared: 0.00011365569809163034

```
print("Regression Coefficients:")
print(model.params)
```

Regression Coefficients:
const 785.098535
Age 0.470893
dtype: float64

```
print("P-values:")
print(model.pvalues)
```

P-values:
const 5.267571e-65
Age 6.829722e-01
dtype: float64

```
print("Confidence Intervals:")
print(model.conf_int())
```

Confidence Intervals:
const 699.089362 871.107708
Age -1.790341 2.732128

```
print(model.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          DailyRate   R-squared: 
0.000
Model:                  OLS   Adj.   R-squared:   F-
-0.001
Method:                 Least Squares   statistic: 
0.1669
Date:                   Mon,31   Aug 2026   Prob (F-statistic): Log-
0.683
Time:                   08:31:37   Likelihood: AIC: 
-10906.
No. Observations:      1470   BIC: 
2.182e+04
Df Residuals:          1468
2.183e+04
Df Model:               1

CovarianceType:         nonrobust
```



```

=====
              coef      std err          t      P>|t|      [0.025
0.975] -----
const
871.108
Age          785.0985      43.847      17.905      0.000      699.089
2.732
===== -0.4709 -1.153 -0.408 -0.683 -1.790 =====
=====
Omnibus:
2.082
Prob(Omnibus):          1174.453   Durbin-Watson:
88.696
Skew:          0.000   Jarque-Bera   (JB):
5.49e-20
Kurtosis:          -0.004   Prob(JB):
159.
===== -1.797 -Cond. No. =====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

print("STATISTICAL ANALYSIS SUMMARY")
print("=" * 40)

```

```

print("\n1. T-TEST")
print("T-statistic:", t_stat)
print("P-value:", p_value)

```

```

if p_value < 0.05:

```

```

    print("Result: Significant difference in Age.")
else:
    print("Result: No significant difference in Age.")

```

```

print("\n2. ANOVA")
print("F-statistic:", f_stat)
print("P-value:", p_value_anova)

```

```

if p_value_anova < 0.05:
    print("Result: Significant difference among Departments.")
else:
    print("Result: No significant difference among Departments.")

```

```

print("\n3. LINEAR REGRESSION")
print("R-squared:", model.rsquared)

```

```
print("Age coefficient:", model.params["Age"])  
print("Age p-value:", model.pvalues["Age"])
```

STATISTICAL ANALYSIS SUMMARY

=====

1. T-TEST

T-statistic: -5.82801185398895

P-value: 1.3797600649439802e-08

Result: Significant difference in Age.

2. ANOVA

F-statistic: 0.7655031924976907

P-value: 0.4652855299935297

Result: No significant difference among Departments.

3. LINEAR REGRESSION

R-squared: 0.00011365569809163034

Age coefficient: 0.47089343162226616

Age p-value: 0.6829722453880929